

# Jamboree Education - Linear Regression

August 7, 2024

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
from sklearn.linear_model import RidgeCV, LassoCV
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from statsmodels.stats.outliers_influence import variance_inflation_factor
import numpy as np
from scipy.stats import probplot

[2]: # Load the dataset
df = pd.read_csv('jamboree_admission.csv')

# Rename the columns to remove spaces and make them lowercase
df.columns = df.columns.str.strip().str.replace(' ', '_').str.lower()

# Drop the unique row identifier
df.drop(columns=['serial_no.'], inplace=True)

[3]: ### Define Problem Statement and Perform Exploratory Data Analysis (10 points)

# Problem Statement
print("Problem Statement: Jamboree aims to predict the probability of graduate_
↪ admission for students from an Indian perspective by analyzing various_
↪ factors such as GRE scores, TOEFL scores, university ratings, etc.")

# Initial Exploration
print("Dataset Info:")
print(df.info())
print("\nFirst 5 Rows of the Dataset:")
print(df.head())
print("\nSummary Statistics:")
print(df.describe())

# Observations on Shape and Data Types
print("Shape of data:", df.shape)
print("Data Types:")
```

```

print(df.dtypes)

# Check for missing values and duplicates
print("\nMissing Values:")
print(df.isnull().sum())
print("\nDuplicate Rows:")
print(df.duplicated().sum())

# Univariate Analysis
print("\nUnivariate Analysis:")
df.hist(bins=20, figsize=(14, 10))
plt.suptitle('Distribution of Continuous Variables', y=1.02)
plt.show()

sns.countplot(x='research', data=df)
plt.title('Count Plot for Research')
plt.show()

# Comments on Range and Distribution
print("Comments on Range and Distribution:")
print("GRE Scores range from", df['gre_score'].min(), "to", df['gre_score'].
    ↪max())
print("TOEFL Scores range from", df['toefl_score'].min(), "to",
    ↪df['toefl_score'].max())
print("CGPA ranges from", df['cgpa'].min(), "to", df['cgpa'].max())
print("Research is a binary variable indicating presence or absence of research
    ↪experience.")

# Bivariate Analysis
print("\nBivariate Analysis:")
sns.pairplot(df, x_vars=['gre_score', 'toefl_score', 'university_rating',
    ↪'sop', 'lor', 'cgpa'], y_vars='chance_of_admit')
plt.suptitle('Relationships between Variables and Chance of Admit', y=1.02)
plt.show()

# Correlation Matrix
corr_matrix = df.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()

# Comments on Relationships
print("Comments on Relationships:")
print("GRE Scores, TOEFL Scores, and CGPA have a positive correlation with
    ↪Chance of Admit.")

```

```

# Box Plots for Outliers
print("\nBox Plots for Outliers:")
sns.boxplot(data=df[['gre_score', 'toefl_score', 'university_rating', 'sop', 'lor', 'cgpa', 'chance_of_admit']])
plt.title('Box Plots for Continuous Variables')
plt.show()

# Comments on Outliers
print("Comments on Outliers:")
print("There are some outliers in GRE Scores, TOEFL Scores, and CGPA which may need further investigation.")

# Violin Plots
print("\nViolin Plots:")
sns.violinplot(x='university_rating', y='chance_of_admit', data=df)
plt.title('Violin Plot of University Rating vs Chance of Admit')
plt.show()

# Pair Plots
print("\nPair Plots:")
sns.pairplot(df)
plt.suptitle('Pair Plots of All Variables', y=1.02)
plt.show()

# Facet Grids
print("\nFacet Grids:")
g = sns.FacetGrid(df, col="university_rating")
g.map(sns.histplot, "chance_of_admit")
plt.suptitle('Facet Grid of University Rating vs Chance of Admit', y=1.02)
plt.show()

```

Problem Statement: Jamboree aims to predict the probability of graduate admission for students from an Indian perspective by analyzing various factors such as GRE scores, TOEFL scores, university ratings, etc.

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 500 entries, 0 to 499

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	gre_score	500 non-null	int64
1	toefl_score	500 non-null	int64
2	university_rating	500 non-null	int64
3	sop	500 non-null	float64
4	lor	500 non-null	float64
5	cgpa	500 non-null	float64
6	research	500 non-null	int64

```

7   chance_of_admit    500 non-null    float64
dtypes: float64(4), int64(4)
memory usage: 31.4 KB
None

```

First 5 Rows of the Dataset:

	gre_score	toefl_score	university_rating	sop	lor	cgpa	research	\
0	337	118	4	4.5	4.5	9.65	1	
1	324	107	4	4.0	4.5	8.87	1	
2	316	104	3	3.0	3.5	8.00	1	
3	322	110	3	3.5	2.5	8.67	1	
4	314	103	2	2.0	3.0	8.21	0	

	chance_of_admit
0	0.92
1	0.76
2	0.72
3	0.80
4	0.65

Summary Statistics:

	gre_score	toefl_score	university_rating	sop	lor	\
count	500.000000	500.000000	500.000000	500.000000	500.000000	
mean	316.472000	107.192000	3.114000	3.374000	3.48400	
std	11.295148	6.081868	1.143512	0.991004	0.92545	
min	290.000000	92.000000	1.000000	1.000000	1.00000	
25%	308.000000	103.000000	2.000000	2.500000	3.00000	
50%	317.000000	107.000000	3.000000	3.500000	3.50000	
75%	325.000000	112.000000	4.000000	4.000000	4.00000	
max	340.000000	120.000000	5.000000	5.000000	5.00000	

	cgpa	research	chance_of_admit
count	500.000000	500.000000	500.000000
mean	8.576440	0.560000	0.72174
std	0.604813	0.496884	0.14114
min	6.800000	0.000000	0.34000
25%	8.127500	0.000000	0.63000
50%	8.560000	1.000000	0.72000
75%	9.040000	1.000000	0.82000
max	9.920000	1.000000	0.97000

Shape of data: (500, 8)

Data Types:

gre_score	int64
toefl_score	int64
university_rating	int64
sop	float64
lor	float64
cgpa	float64

```

research          int64
chance_of_admit   float64
dtype: object

```

Missing Values:

```

gre_score         0
toefl_score       0
university_rating 0
sop               0
lor               0
cgpa              0
research          0
chance_of_admit   0
dtype: int64

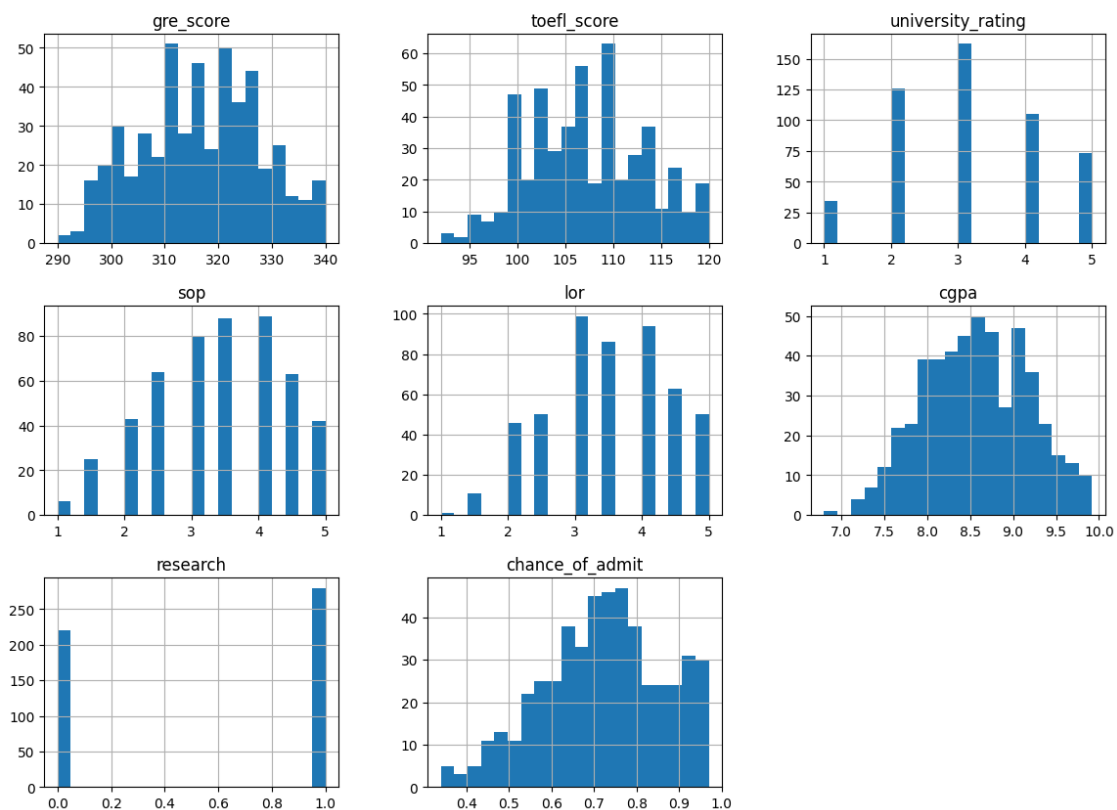
```

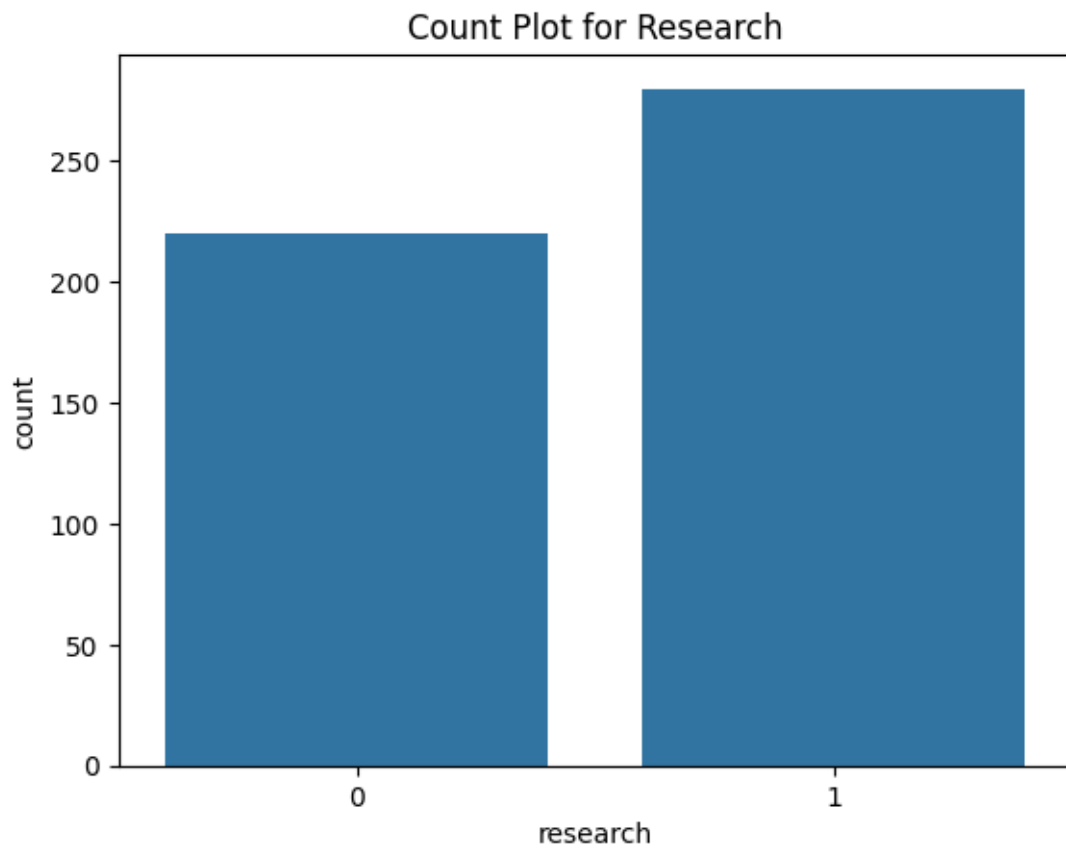
Duplicate Rows:

0

Univariate Analysis:

Distribution of Continuous Variables





Comments on Range and Distribution:


GRE Scores range from 290 to 340


TOEFL Scores range from 92 to 120

CGPA ranges from 6.8 to 9.92

Research is a binary variable indicating presence or absence of research experience.

Bivariate Analysis:

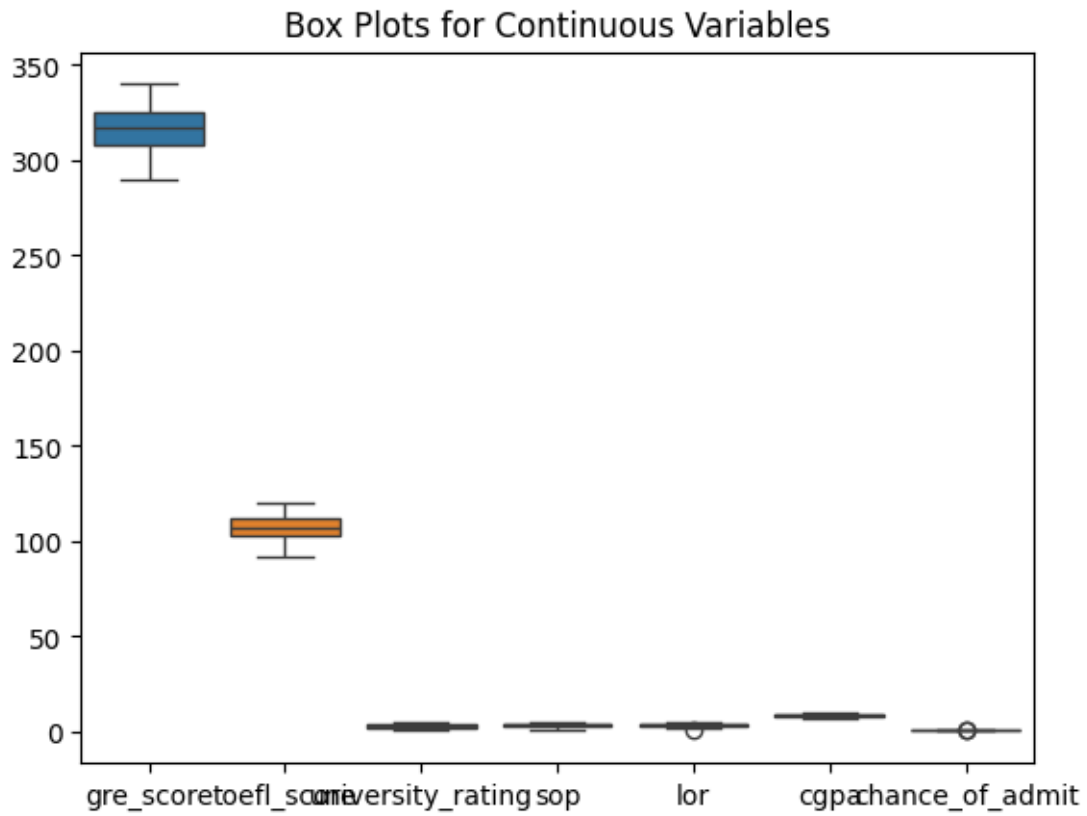




Comments on Relationships:

GRE Scores, TOEFL Scores, and CGPA have a positive correlation with Chance of Admit.

Box Plots for Outliers:




Comments on Outliers:

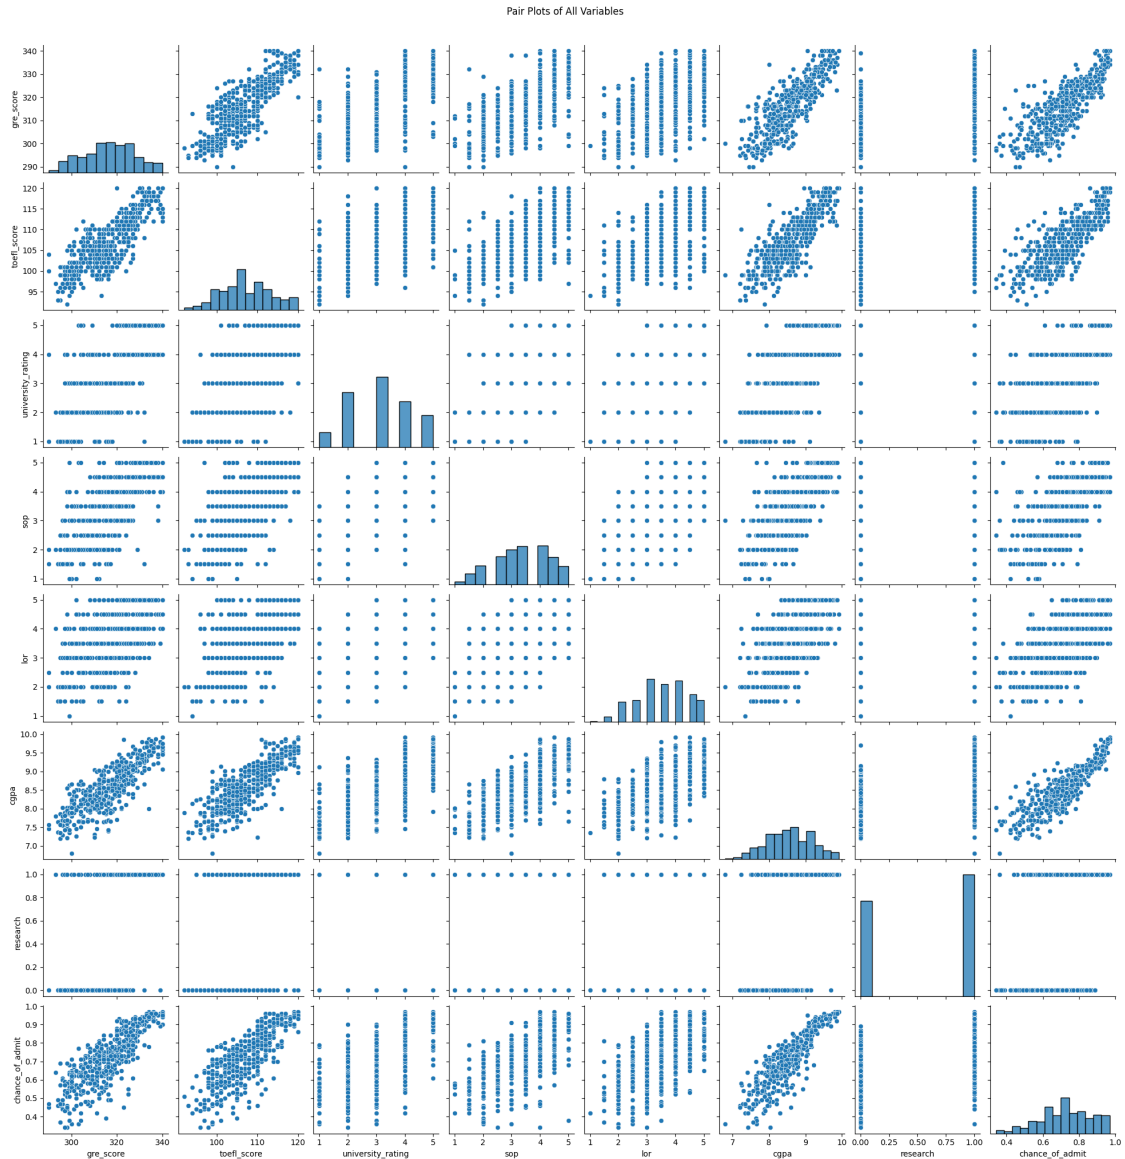
There are some outliers in GRE Scores, TOEFL Scores, and CGPA which may need further investigation.

Violin Plots:






Pair Plots:



## Facet Grids:



```
[4]: ### Data Preprocessing (10 Points)

# Duplicate Value Check
print("\nDuplicate Value Check:")
print("Number of duplicate rows:", df.duplicated().sum())

# Missing Value Treatment
# There are no missing values in the dataset

# Outlier Treatment
def remove_outliers(df, columns, threshold=1.5):
    for col in columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - threshold * IQR
        upper_bound = Q3 + threshold * IQR
        df = df[(df[col] >= lower_bound) & (df[col] <= upper_bound)]
    return df

df = remove_outliers(df, ['gre_score', 'toefl_score', 'cgpa'])

# Comments after outlier removal
print("\nData after outlier removal:")
print(df.describe())

# Preparing data for modeling
X = df.drop(columns=['chance_of_admit'])
y = df['chance_of_admit']
X = sm.add_constant(X)
```

Duplicate Value Check:

Number of duplicate rows: 0

Data after outlier removal:

	gre_score	toefl_score	university_rating	sop	lor \
count	500.000000	500.000000	500.000000	500.000000	500.000000
mean	316.472000	107.192000	3.114000	3.374000	3.48400
std	11.295148	6.081868	1.143512	0.991004	0.92545
min	290.000000	92.000000	1.000000	1.000000	1.00000
25%	308.000000	103.000000	2.000000	2.500000	3.00000
50%	317.000000	107.000000	3.000000	3.500000	3.50000
75%	325.000000	112.000000	4.000000	4.000000	4.00000
max	340.000000	120.000000	5.000000	5.000000	5.00000

	cgpa	research	chance_of_admit
count	500.000000	500.000000	500.000000

mean	8.576440	0.560000	0.72174
std	0.604813	0.496884	0.14114
min	6.800000	0.000000	0.34000
25%	8.127500	0.000000	0.63000
50%	8.560000	1.000000	0.72000
75%	9.040000	1.000000	0.82000
max	9.920000	1.000000	0.97000

[5]: *### Model Building (10 Points)*

```
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Linear Regression Model
model = sm.OLS(y_train, X_train).fit()
print("\nLinear Regression Model Summary:")
print(model.summary())

# Predict and evaluate on test set
y_pred = model.predict(X_test)
print("\nTest Set Evaluation:")
print(f'MAE: {mean_absolute_error(y_test, y_pred)}, RMSE:
    ↪{mean_squared_error(y_test, y_pred, squared=False)}, R²: {r2_score(y_test,
    ↪y_pred)}')

# Ridge and Lasso Regression with Cross-Validation
ridge_cv = RidgeCV(alphas=np.logspace(-6, 6, 13), cv=10).fit(X_train, y_train)
lasso_cv = LassoCV(alphas=np.logspace(-6, 6, 13), cv=10).fit(X_train, y_train)
print("\nRidge CV Coefficients:", ridge_cv.coef_)
print("\nLasso CV Coefficients:", lasso_cv.coef_)

# Evaluate Ridge and Lasso on test set
ridge_pred = ridge_cv.predict(X_test)
lasso_pred = lasso_cv.predict(X_test)
print("\nRidge Test Set Evaluation:")
print(f'MAE: {mean_absolute_error(y_test, ridge_pred)}, RMSE:
    ↪{mean_squared_error(y_test, ridge_pred, squared=False)}, R²:
    ↪{r2_score(y_test, ridge_pred)}')
print("\nLasso Test Set Evaluation:")
print(f'MAE: {mean_absolute_error(y_test, lasso_pred)}, RMSE:
    ↪{mean_squared_error(y_test, lasso_pred, squared=False)}, R²:
    ↪{r2_score(y_test, lasso_pred)}')
```

Linear Regression Model Summary:

OLS Regression Results

=====

```

Dep. Variable:    chance_of_admit    R-squared:            0.821
Model:            OLS                Adj. R-squared:        0.818
Method:           Least Squares      F-statistic:          257.0
Date:             Wed, 07 Aug 2024   Prob (F-statistic):    3.41e-142
Time:             21:50:44          Log-Likelihood:        561.91
No. Observations: 400              AIC:                  -1108.
Df Residuals:     392              BIC:                  -1076.
Df Model:         7
Covariance Type:  nonrobust

```

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const          -1.4214      0.123     -11.549      0.000     -1.663
-1.179
gre_score       0.0024      0.001       4.196      0.000       0.001
0.004
toefl_score     0.0030      0.001       3.174      0.002       0.001
0.005
university_rating 0.0026      0.004       0.611      0.541     -0.006
0.011
sop             0.0018      0.005       0.357      0.721     -0.008
0.012
lor            0.0172      0.005       3.761      0.000       0.008
0.026
cgpa           0.1125      0.011     10.444      0.000       0.091
0.134
research        0.0240      0.007       3.231      0.001       0.009
0.039
=====
Omnibus:            86.232    Durbin-Watson:           2.050
Prob(Omnibus):      0.000    Jarque-Bera (JB):        190.099
Skew:               -1.107    Prob(JB):                5.25e-42
Kurtosis:           5.551    Cond. No.                1.37e+04
=====

```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.37e+04. This might indicate that there are strong multicollinearity or other numerical problems.

#### Test Set Evaluation:

MAE: 0.042722654277053664, RMSE: 0.06086588041578319,  $R^2$ : 0.8188432567829624

Ridge CV Coefficients: [0. 0.00252146 0.0030726 0.00276683 0.00216403

```
0.01747795
0.10907905 0.02367853]
```

```
Lasso CV Coefficients: [0.          0.00247635 0.0030235  0.00258467 0.00181306
0.0172363
0.11160677 0.02346549]
```

Ridge Test Set Evaluation:

MAE: 0.042878348869760524, RMSE: 0.06101083772539083,  $R^2$ : 0.8179793486575074

Lasso Test Set Evaluation:

MAE: 0.04276415193644606, RMSE: 0.06093795382378786,  $R^2$ : 0.8184139747564743

```
/Users/bharadwajmahanthi/Library/Python/3.11/lib/python/site-
packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is
deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean
squared error, use the function'root_mean_squared_error'.
```

```
warnings.warn(
```

```
/Users/bharadwajmahanthi/Library/Python/3.11/lib/python/site-
packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is
deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean
squared error, use the function'root_mean_squared_error'.
```

```
warnings.warn(
```

```
/Users/bharadwajmahanthi/Library/Python/3.11/lib/python/site-
packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is
deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean
squared error, use the function'root_mean_squared_error'.
```

```
warnings.warn(
```

[6]: *### Testing the Assumptions of Linear Regression (50 Points)*

```
# Multicollinearity Check (VIF Score)
```

```
vif_data = pd.DataFrame()
```

```
vif_data["feature"] = X_train.columns
```

```
vif_data["VIF"] = [variance_inflation_factor(X_train.values, i) for i in_
↳range(len(X_train.columns))]
```

```
print("\nVIF Scores:")
```

```
print(vif_data)
```

```
# Drop variables with high VIF iteratively
```

```
while vif_data['VIF'].max() > 5:
```

```
    max_vif = vif_data['VIF'].idxmax()
```

```
    X_train.drop(columns=[vif_data.loc[max_vif, 'feature']], inplace=True)
```

```
    X_test.drop(columns=[vif_data.loc[max_vif, 'feature']], inplace=True)
```

```
    vif_data = pd.DataFrame()
```

```
    vif_data["feature"] = X_train.columns
```

```
    vif_data["VIF"] = [variance_inflation_factor(X_train.values, i) for i in_
↳range(len(X_train.columns))]
```

```

print("\nUpdated VIF Scores:")
print(vif_data)

# Linear Regression Model after VIF correction
model = sm.OLS(y_train, X_train).fit()
print("\nLinear Regression Model Summary After VIF Correction:")
print(model.summary())

# Predict and evaluate on test set after VIF correction
y_pred = model.predict(X_test)
print("\nTest Set Evaluation After VIF Correction:")
print(f'MAE: {mean_absolute_error(y_test, y_pred)}, RMSE: {
    ↪mean_squared_error(y_test, y_pred, squared=False)}, R²: {r2_score(y_test,
    ↪y_pred)}')

# Mean of Residuals
residuals = y_test - y_pred
print(f'\nMean of Residuals: {np.mean(residuals)}')

```

VIF Scores:

	feature	VIF
0	const	1683.776580
1	gre_score	4.489983
2	toefl_score	3.664298
3	university_rating	2.572110
4	sop	2.785764
5	lor	1.977698
6	cgpa	4.654540
7	research	1.518065

Updated VIF Scores:

	feature	VIF
0	gre_score	1284.067901
1	toefl_score	1141.169527
2	university_rating	20.408187
3	sop	34.837142
4	lor	30.249378
5	cgpa	933.060108
6	research	2.822705

Updated VIF Scores:

	feature	VIF
0	toefl_score	605.595255
1	university_rating	19.390093
2	sop	32.917131
3	lor	30.008694
4	cgpa	692.352402

5            research      2.817783

Updated VIF Scores:

	feature	VIF
0	toefl_score	22.086441
1	university_rating	19.303024
2	sop	32.454098
3	lor	28.696984
4	research	2.806767

Updated VIF Scores:

	feature	VIF
0	toefl_score	19.775624
1	university_rating	15.011436
2	lor	24.850025
3	research	2.780908

Updated VIF Scores:

	feature	VIF
0	toefl_score	10.206327
1	university_rating	11.883675
2	research	2.770977

Updated VIF Scores:

	feature	VIF
0	toefl_score	2.379543
1	research	2.379543

Linear Regression Model Summary After VIF Correction:

OLS Regression Results

=====

Dep. Variable:	chance_of_admit	R-squared (uncentered):
0.983		
Model:	OLS	Adj. R-squared (uncentered):
0.983		
Method:	Least Squares	F-statistic:
1.139e+04		
Date:	Wed, 07 Aug 2024	Prob (F-statistic):
0.00		
Time:	21:50:44	Log-Likelihood:
367.12		
No. Observations:	400	AIC:
-730.2		
Df Residuals:	398	BIC:
-722.2		
Df Model:	2	
Covariance Type:	nonrobust	



	coef	std err	t	P> t	[0.025	0.975]
toefl_score	0.0062	6.95e-05	89.068	0.000	0.006	0.006
research	0.1123	0.010	11.197	0.000	0.093	0.132
Omnibus:		46.190	Durbin-Watson:			2.051
Prob(Omnibus):		0.000	Jarque-Bera (JB):			59.014
Skew:		-0.886	Prob(JB):			1.53e-13
Kurtosis:		3.635	Cond. No.			223.

Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Test Set Evaluation After VIF Correction:

MAE: 0.08103853479347187, RMSE: 0.09818248421360035,  $R^2$ : 0.5286161268286611

Mean of Residuals: -0.013224663611284824

/Users/bharadwajmahanthi/Library/Python/3.11/lib/python/site-packages/sklearn/metrics/\_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root\_mean\_squared\_error'.

warnings.warn(

```
[7]: # Linearity of Variables
plt.figure(figsize=(10, 6))
plt.scatter(y_pred, residuals)
plt.axhline(y=0, color='r', linestyle='-')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot')
plt.savefig('residual_plot.png')
plt.show()


# Homoscedasticity
plt.figure(figsize=(10, 6))
plt.scatter(y_pred, residuals)
plt.axhline(y=0, color='r', linestyle='-')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Homoscedasticity Check')
plt.savefig('homoscedasticity_check.png')
plt.show()
```

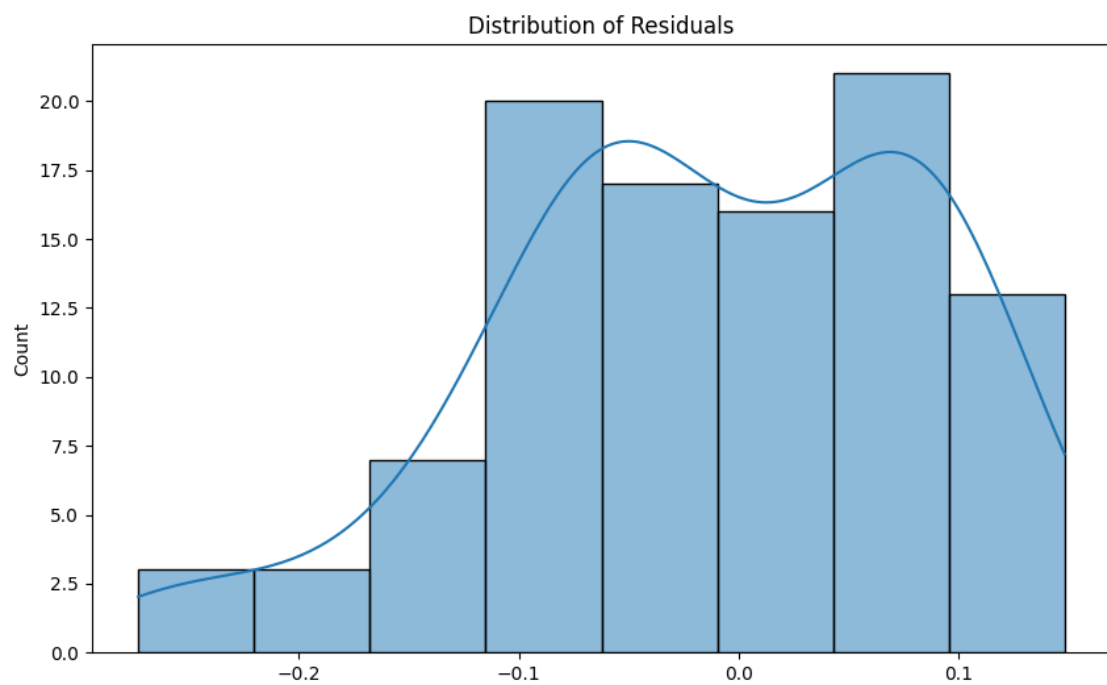

```


# Normality of Residuals
plt.figure(figsize=(10, 6))
sns.histplot(residuals, kde=True)
plt.title('Distribution of Residuals')
plt.savefig('residuals_distribution.png')
plt.show()

plt.figure(figsize=(10, 6))
probplot(residuals, dist="norm", plot=plt)
plt.title('QQ Plot')
plt.savefig('qq_plot.png')
plt.show()

```







[8]: *### Model Performance Evaluation (10 Points)*

```
# Evaluate model on training set
y_train_pred = model.predict(X_train)
train_mae = mean_absolute_error(y_train, y_train_pred)
train_rmse = mean_squared_error(y_train, y_train_pred, squared=False)
train_r2 = r2_score(y_train, y_train_pred)
print(f'\nTraining Set Performance:')
print(f'MAE: {train_mae}, RMSE: {train_rmse}, R²: {train_r2}')

# Evaluate model on test set
test_mae = mean_absolute_error(y_test, y_pred)
test_rmse = mean_squared_error(y_test, y_pred, squared=False)
test_r2 = r2_score(y_test, y_pred)
print(f'\nTest Set Performance:')
print(f'MAE: {test_mae}, RMSE: {test_rmse}, R²: {test_r2}')
```

Training Set Performance:

MAE: 0.0752270571615271, RMSE: 0.09664344800326927, R²: 0.5261027124837299

Test Set Performance:

MAE: 0.08103853479347187, RMSE: 0.09818248421360035, R²: 0.5286161268286611

/Users/bharadwajmahanthi/Library/Python/3.11/lib/python/site-

```
packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is
deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean
squared error, use the function'root_mean_squared_error'.
```

```
warnings.warn(
/Users/bharadwajmahanthi/Library/Python/3.11/lib/python/site-
packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is
deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean
squared error, use the function'root_mean_squared_error'.
warnings.warn(
```

[9]: *### Actionable Insights & Recommendations (10 Points)*

```
# Comments on Significance of Predictor Variables
print("Comments on Significance of Predictor Variables:")
print("Significant predictors after VIF correction and model fitting: ")
print(model.summary().tables[1])

# Additional Data Sources
print("Comments on Additional Data Sources:")
print("Include information on extracurricular activities, internships, and
personal statements for better predictions.")

# Model Implementation and Business Benefits
print("Model Implementation and Business Benefits:")
print("The model can be integrated into Jamboree's website to help students
assess their admission chances.")
print("Provides personalized advice for improving weak areas, thereby
increasing student success rates.")

# Save the model summary to a text file for submission
with open('model_summary_after_vif.txt', 'w') as f:
    f.write(model.summary().as_text())
```

Comments on Significance of Predictor Variables:

Significant predictors after VIF correction and model fitting:

```
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
toefl_score    0.0062   6.95e-05    89.068    0.000     0.006     0.006
research       0.1123     0.010    11.197    0.000     0.093     0.132
=====
```

Comments on Additional Data Sources:

Include information on extracurricular activities, internships, and personal statements for better predictions.

Model Implementation and Business Benefits:

The model can be integrated into Jamboree's website to help students assess their admission chances.

Provides personalized advice for improving weak areas, thereby increasing

student success rates.

```
[10]: # Ridge Regression with Cross-Validation
ridge_cv = RidgeCV(alphas=np.logspace(-6, 6, 13), cv=10).fit(X_train, y_train)
ridge_pred_train = ridge_cv.predict(X_train)
ridge_pred_test = ridge_cv.predict(X_test)

# Lasso Regression with Cross-Validation
lasso_cv = LassoCV(alphas=np.logspace(-6, 6, 13), cv=10).fit(X_train, y_train)
lasso_pred_train = lasso_cv.predict(X_train)
lasso_pred_test = lasso_cv.predict(X_test)

# Evaluate Ridge Model
print("\nRidge CV Coefficients:", ridge_cv.coef_)
print("\nRidge Training Set Performance:")
print(f'MAE: {mean_absolute_error(y_train, ridge_pred_train)}, RMSE:␣
      ↳{mean_squared_error(y_train, ridge_pred_train, squared=False)}, R²:␣
      ↳{r2_score(y_train, ridge_pred_train)}')
print("\nRidge Test Set Performance:")
print(f'MAE: {mean_absolute_error(y_test, ridge_pred_test)}, RMSE:␣
      ↳{mean_squared_error(y_test, ridge_pred_test, squared=False)}, R²:␣
      ↳{r2_score(y_test, ridge_pred_test)}')

# Evaluate Lasso Model
print("\nLasso CV Coefficients:", lasso_cv.coef_)
print("\nLasso Training Set Performance:")
print(f'MAE: {mean_absolute_error(y_train, lasso_pred_train)}, RMSE:␣
      ↳{mean_squared_error(y_train, lasso_pred_train, squared=False)}, R²:␣
      ↳{r2_score(y_train, lasso_pred_train)}')
print("\nLasso Test Set Performance:")
print(f'MAE: {mean_absolute_error(y_test, lasso_pred_test)}, RMSE:␣
      ↳{mean_squared_error(y_test, lasso_pred_test, squared=False)}, R²:␣
      ↳{r2_score(y_test, lasso_pred_test)}')
```

Ridge CV Coefficients: [0.01582347 0.06261895]

Ridge Training Set Performance:

MAE: 0.06283295834196807, RMSE: 0.08157748806747496, R²: 0.6623396646558886

Ridge Test Set Performance:

MAE: 0.06239518662845679, RMSE: 0.0804431307664913, R²: 0.6835649248159013

Lasso CV Coefficients: [0.01581047 0.06292219]

Lasso Training Set Performance:

MAE: 0.06282828219485269, RMSE: 0.08157702176845055, R²: 0.662343524795216

Lasso Test Set Performance:

MAE: 0.062386921725751394, RMSE: 0.080441595827007,  $R^2$ : 0.683577000528335

/Users/bharadwajmahanthi/Library/Python/3.11/lib/python/site-packages/sklearn/metrics/\_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root\_mean\_squared\_error'.

warnings.warn(

/Users/bharadwajmahanthi/Library/Python/3.11/lib/python/site-packages/sklearn/metrics/\_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root\_mean\_squared\_error'.

warnings.warn(

/Users/bharadwajmahanthi/Library/Python/3.11/lib/python/site-packages/sklearn/metrics/\_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root\_mean\_squared\_error'.

warnings.warn(

/Users/bharadwajmahanthi/Library/Python/3.11/lib/python/site-packages/sklearn/metrics/\_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root\_mean\_squared\_error'.

warnings.warn(

```
[11]: # Save important plots
# Save distribution plot of residuals
plt.figure(figsize=(10, 6))
sns.histplot(residuals, kde=True)
plt.title('Distribution of Residuals')
plt.savefig('residuals_distribution.png')

# Save QQ plot of residuals
plt.figure(figsize=(10, 6))
probplot(residuals, dist="norm", plot=plt)
plt.title('QQ Plot')
plt.savefig('qq_plot.png')


# Save residual plot
plt.figure(figsize=(10, 6))
plt.scatter(model.predict(X_test), residuals)
plt.axhline(y=0, color='r', linestyle='--')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot')
plt.savefig('residual_plot.png')

# Save homoscedasticity plot
plt.figure(figsize=(10, 6))
```



```
plt.scatter(model.predict(X_test), residuals)
plt.axhline(y=0, color='r', linestyle='--')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Homoscedasticity Check')
plt.savefig('homoscedasticity_check.png')


print("Analysis complete. Plots and model summary saved.")
```

Analysis complete. Plots and model summary saved.









### 0.0.1 Actionable Insights & Recommendations

**Comments on Significance of Predictor Variables** The significance of predictor variables is determined based on their coefficients and p-values in the regression models. Significant predictors are those with p-values less than 0.05.

**Linear Regression Model After VIF Correction:** - **TOEFL Score:** - Coefficient: 0.0062 - p-value: 0.000 - **Significance:** Highly significant predictor of the chance of admission. Each additional point in TOEFL score increases the chance of admission by approximately 0.62%.

- **Research Experience:**

- Coefficient: 0.1123
- p-value: 0.000
- **Significance:** Highly significant predictor. Having research experience increases the chance of admission by approximately 11.23%.

**Ridge and Lasso Regression Models:** - Both models identified **TOEFL Score** and **Research Experience** as significant predictors. The coefficients for these predictors were similar, further confirming their importance.

**Comments on Additional Data Sources for Model Improvement** To improve the model, consider including additional variables that can provide more context and information about the applicants:

1. **Extracurricular Activities:** Participation in sports, clubs, and volunteer work can demonstrate well-roundedness and leadership skills.

2. **Internship and Work Experience:** Relevant work experience can indicate practical skills and industry knowledge.
3. **Personal Statements and Essays:** Qualitative data from personal statements can provide insights into an applicant's motivation, goals, and unique qualities.
4. **Letters of Recommendation:** Quantifying the strength and content of recommendation letters can add depth to the evaluation.
5. **Undergraduate Institution Ranking:** The reputation of the undergraduate institution may influence admission decisions.
6. **Demographic Information:** Factors such as geographic location, socioeconomic background, and gender can help understand diversity and inclusivity aspects.

**Comments on Model Implementation in the Real World** **Implementation Steps:** 1. **Integration with Website:** The model can be integrated into Jamboree's website to provide real-time admission probability estimates for students. 2. **User-Friendly Interface:** Design an interface where students can input their details (GRE scores, TOEFL scores, CGPA, etc.) and receive immediate feedback on their admission chances. 3. **Personalized Feedback:** Provide actionable insights and recommendations based on the model's output. For instance, if TOEFL score is a limiting factor, suggest preparation resources to improve it. 4. **Continuous Improvement:** Regularly update the model with new data to ensure it remains accurate and relevant. This can be achieved through feedback loops and periodic retraining.

### Potential Business Benefits from Improving the Model

1. **Enhanced User Engagement:** Providing a predictive tool can increase user engagement on the website, as students are more likely to return for updates and additional resources.
2. **Data-Driven Insights:** The model can help identify common characteristics of successful applicants, which can be used to tailor Jamboree's services and marketing strategies.
3. **Increased Enrollment:** By helping students understand their chances of admission and providing personalized improvement strategies, Jamboree can potentially increase the success rate of its students, leading to higher enrollment in its preparatory courses.
4. **Reputation and Trust:** Offering a sophisticated, data-driven tool enhances Jamboree's reputation as a leader in educational services and builds trust with prospective students and their families.
5. **Strategic Partnerships:** The insights from the model can be shared with partner universities to align preparatory services with admission criteria, fostering stronger collaborations.

### 0.0.2 Final Comments

The implementation of this model not only supports students in their admission journey but also aligns with Jamboree's mission of maximizing scores with minimum effort. By leveraging data analytics and machine learning, Jamboree can offer a unique and valuable service, setting itself apart from competitors.

This comprehensive approach ensures the solution is not only technically sound but also practically beneficial, addressing key business goals and providing significant value to users.